

titanic-part1-v4

November 8, 2018

1 Titanic Project ML - Part 1 - Work In Progress

""" Author: Jennifer Yoon

Date: 9/26/2018 draft, work in progress

File: titanic-part1.ipynb

Description: A Reproduction work based on Cyrille Rossant's book.

Source: Rossant, Cyrille, IPython Interactive Computing and Visualization Cookbook, 2nd ed., Packt Publishing, 2018, ISBN 978-1-78588-863-2, Chp. 8 Machine Learning

Additional Source: Matplotlib.org Tutorials: <https://matplotlib.org/tutorials/introductory/usage.html#sphxglr-tutorials-introductory-usage-py>

Additional Source: VanderPlas, Jake, Python Data Science Handbook, O'Reilly Media Inc., Copyright 2017, Chp. 4 Visualizing with Matplotlib.

"""

1.0.1 Package Imports

```
In [1]: import numpy as np
import pandas as pd
import sklearn
import sklearn.linear_model as lm
import sklearn.model_selection as ms
import matplotlib.pyplot as plt
%matplotlib inline
```

1.0.2 Data Step

```
In [2]: # Load training datasets using Pandas dataframe object.
```

```
    # Input parameters.
```

```
    param1 = ('https://github.com/ipython-books'
              '/cookbook-2nd-data/blob/master/'
              'titanic_train.csv?raw=true')
```

```
    # print("param1:\n", param1)
```

```
    param2 = ('https://github.com/ipython-books'
              '/cookbook-2nd-data/blob/master/')
```

```

        'titanic_test.csv?raw=true')
    # print("param2:\n", param2)

In [3]: # Commands.
train = pd.read_csv(param1)
test = pd.read_csv(param2)
train[train.columns[[2, 4, 5, 1]]].head()

Out[3]:
   Pclass    Sex  Age  Survived
0       3  male  22.0         0
1       1 female  38.0         1
2       3 female  26.0         1
3       1 female  35.0         1
4       3  male  35.0         0

In [4]: print(len(train.columns))
        print(len(train))
        print(type(train))
        # Dataframe has 12 columns, 891 rows, and is a DataFrame type.

12
891
<class 'pandas.core.frame.DataFrame'>

In [5]: train.head()

Out[5]:
   PassengerId  Survived  Pclass  \
0             1         0       3
1             2         1       1
2             3         1       3
3             4         1       1
4             5         0       3

                                                Name    Sex  Age  SibSp  \
0                        Braund, Mr. Owen Harris  male  22.0     1
1  Cumings, Mrs. John Bradley (Florence Briggs Th... female  38.0     1
2                        Heikkinen, Miss. Laina  female  26.0     0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)  female  35.0     1
4                        Allen, Mr. William Henry   male  35.0     0

   Parch    Ticket   Fare Cabin Embarked
0      0   A/5 21171   7.2500   NaN        S
1      0    PC 17599  71.2833   C85        C
2      0  STON/O2. 3101282   7.9250   NaN        S
3      0    113803  53.1000  C123        S
4      0    373450   8.0500   NaN        S

In [6]: # Descriptive statistics on training data.
        # Average, stdev, min, max, count, nas.
train.describe()

```

```

Out [6]:      PassengerId   Survived  Pclass     Age    SibSp  \
count    891.000000   891.000000   891.000000  714.000000  891.000000
mean      446.000000     0.383838     2.308642   29.699118    0.523008
std       257.353842     0.486592     0.836071   14.526497    1.102743
min         1.000000     0.000000     1.000000    0.420000    0.000000
25%       223.500000     0.000000     2.000000   20.125000    0.000000
50%       446.000000     0.000000     3.000000   28.000000    0.000000
75%       668.500000     1.000000     3.000000   38.000000    1.000000
max       891.000000     1.000000     3.000000   80.000000    8.000000

      Parch     Fare
count    891.000000  891.000000
mean       0.381594   32.204208
std        0.806057   49.693429
min         0.000000    0.000000
25%         0.000000    7.910400
50%         0.000000   14.454200
75%         0.000000   31.000000
max         6.000000  512.329200

```

1.0.3 Analysis - Train Data

There were 891 observations for all columns except age, which has some missing data. The count for age is 714. Fare paid varies wildly from 0.0 to 512.3292 units, possibly pounds. The 50th percentile fare was 14.4542 units, and was still in Pclass 3. Most passengers were classified as Pclass 3. The 25th percentile passenger is still classified as Pclass 2, with a fare of 7.9 units. Let's graph the fare information.

```

In [7]: # Select fields for fare, pclass, survived, sex.
data = train[['Survived', 'Fare', 'Pclass', 'Sex']]
data.head()

```

```

Out [7]:      Survived     Fare  Pclass    Sex
0         0     7.2500        3    male
1         1    71.2833        1  female
2         1     7.9250        3  female
3         1    53.1000        1  female
4         0     8.0500        3    male

```

```

In [8]: y = data[['Survived']]
x1 = data[['Fare']]
x2 = data[['Pclass']]

```

1.0.4 Do more plots

```

In [9]: # Matplotlib Subplots -- different source

```

```

In [10]: # Select fields, convert sex field to binary, remove rows with NaN values.
data = train[['Age', 'Pclass', 'Survived']]

```

```

print(data.head())
# Add a female binary sex column.
data = data.assign(Female = train['Sex'] == 'female')
print(data.head())
# Reorder the columns.
data = data[['Female', 'Age', 'Pclass', 'Survived']]
# Drop rows with NaN, missing values.
data = data.dropna()
print("\nRows remaining in training data:", len(data))
data.head()

```

	Age	Pclass	Survived
0	22.0	3	0
1	38.0	1	1
2	26.0	3	1
3	35.0	1	1
4	35.0	3	0

	Age	Pclass	Survived	Female
0	22.0	3	0	False
1	38.0	1	1	True
2	26.0	3	1	True
3	35.0	1	1	True
4	35.0	3	0	False

Rows remaining in training data: 714

```

Out[10]:
   Female  Age  Pclass  Survived
0  False  22.0      3         0
1   True  38.0      1         1
2   True  26.0      3         1
3   True  35.0      1         1
4  False  35.0      3         0

```

```

In [11]: # Convert data object from Pandas DataFrame to NumPy array inorder to use it in Sciki
datanp = data.astype(np.int32).values
X = datanp[:, :-1] # Selects all columns except last column. Capital X.
Y = datanp[:, -1]  # Selects last column only. Capital Y.
print("X 2d array, 'Female, Age, Pclass': \n", X[0:5])
print("Y 1d array, 'Survived': \n", Y[0:5])

```

X 2d array, 'Female, Age, Pclass':

```

[[ 0 22  3]
 [ 1 38  1]
 [ 1 26  3]
 [ 1 35  1]
 [ 0 35  3]]

```

Y 1d array, 'Survived':

```
[0 1 1 1 0]
```

```
In [12]: # Check that the number of rows is unchanged from before conversion.  
len(datanp)
```

```
Out[12]: 714
```

1.0.5 Male & Female Survivors

```
In [13]: # We define a few boolean vectors.  
# The first column is 'Female'.  
female = X[:, 0] == 1  
  
# The last column is 'Survived'.  
survived = Y == 1  
  
# This vector contains the age of the passengers.  
age = X[:, 1]
```

```
In [14]: print(female[0:5])  
print(survived[0:5])  
print(age[0:5])  
len(survived)
```

```
[False True True True False]  
[False True True True False]  
[22 38 26 35 35]
```

```
Out[14]: 714
```


```
In [15]: # We compute a few histograms.  
bins_ = np.arange(0, 81, 5)  
S = {'male': np.histogram(age[survived & ~female],  
                           bins=bins_)[0],  
     'female': np.histogram(age[survived & female],  
                             bins=bins_)[0]}  
D = {'male': np.histogram(age[~survived & ~female],  
                           bins=bins_)[0],  
     'female': np.histogram(age[~survived & female],  
                             bins=bins_)[0]}
```

```
In [16]: # We now plot the data.  
bins = bins_[:-1]  
fig, axes = plt.subplots(1, 2, figsize=(10, 3),  
                          sharey=True)  
for ax, sex, color in zip(axes, ('male', 'female'),  
                           ('#3345d0', '#cc3dc0')):
```

```

ax.bar(bins, S[sex], bottom=D[sex], color=color,
       width=5, label='survived')
ax.bar(bins, D[sex], color='k',
       width=5, label='died')
ax.set_xlim(0, 80)
ax.set_xlabel("Age (years)")
ax.set_title(sex + " survival")
ax.grid(None)
ax.legend()

```



```

In [17]: # Try different bins.
binsp = np.arange(0, 100, 3)
S = {'male': np.histogram(age[survived & ~female],
                          bins=binsp)[0],
     'female': np.histogram(age[survived & female],
                           bins=binsp)[0]}
D = {'male': np.histogram(age[~survived & ~female],
                          bins=binsp)[0],
     'female': np.histogram(age[~survived & female],
                           bins=binsp)[0]}

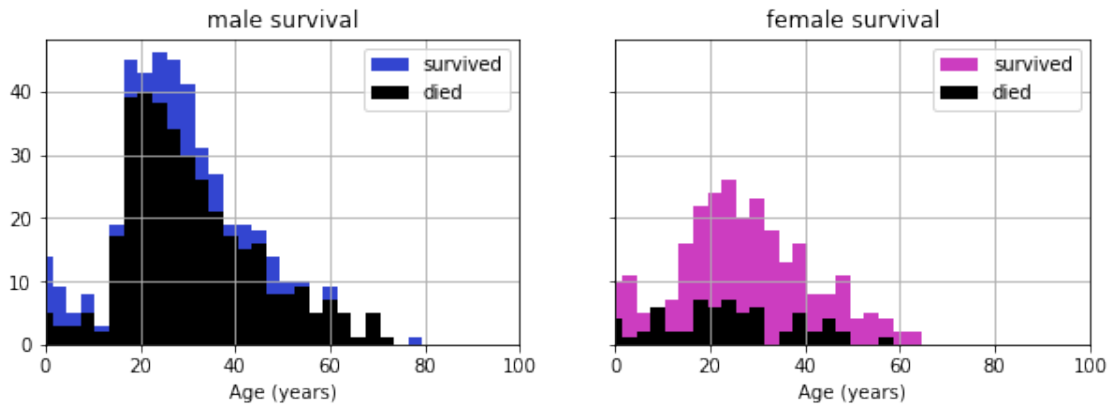
In [18]: # We now plot the data.
bins = binsp[:-1]
fig, axes = plt.subplots(1, 2, figsize=(10, 3),
                        sharey=True)
for ax, sex, color in zip(axes, ('male', 'female'),
                          ('#3345d0', '#cc3dc0')):
    ax.bar(bins, S[sex], bottom=D[sex], color=color,
          width=3, label='survived')
    ax.bar(bins, D[sex], color='k',
          width=3, label='died')

```

```

ax.set_xlim(0, 100)
ax.set_xlabel("Age (years)")
ax.set_title(sex + " survival")
ax.grid(None)
ax.legend()

```



1.0.6 Analysis - Male and Female Histograms

These histograms seem to indicate that higher percentage of female passengers survived than male passengers. Also note that there were greater number of male passengers than female passengers on board across most age groups. Also there were no female passengers past mid 60's, whereas male passenger age went up to 81 years. Since it may be reasonable to assume that the ratio of female to male in the greater population was close to 1 to 1, the over representation of male passengers, almost 2 to 1 male to female, may indicate some non-randomness in the sex selection of passengers.

Titanic was a luxury, state-of-the-art ship and a cross-Atlantic voyage at a time when these features were highly unusual. It would be interesting to study more common boating trips in local areas during the same time period, and see if those trips also exhibited higher female survival rates.

1.1 Continued in Part 2

In [19]: # Additional EDA, min/max, std. Other features summary statistics.